

# PRAVEER NAGARAJA BYNDOOR

(315) 395-5316 | [pbyndoor@syr.edu](mailto:pbyndoor@syr.edu) | [LinkedIn](#) | [GitHub](#)

## EDUCATION

### Syracuse University | School of Information Studies

M.S., Information Systems

Dec 2026

Syracuse, NY

- **GPA:** 4.0/4.0
- **Achievements:** Certificate of Advanced Study: Applied Data Science
- **Coursework:** Applied Machine Learning, Introduction to Data Science, Data Administration and Database Management, Visual Analytic Dashboards, Advanced Big Data Management and Data Warehousing

### Visvesvaraya Technological University | M V J College of Engineering

B.E., Chemical Engineering

Jul 2022

Bengaluru, India

- **GPA:** 3.5/4.0

## PROJECTS

### F1 Applied Machine Learning

- Engineered 12+ predictive features from 10k+ lap records - including rolling lap time averages, tire degradation patterns, and position deltas - improving model  $R^2$  from 0.73 to 0.85
- Trained and compared 4 ensemble models (Random Forest, Extra Trees, Gradient Boosting, CatBoost) using scikit-learn with 5-fold cross-validation, achieving MAE of 2.1 laps on optimal pit-stop predictions
- Tuned hyperparameters via GridSearchCV across 48 combinations, selecting best model based on  $R^2$ , MAE, and RMSE
- Identified top 3 performance drivers via feature importance analysis accounting for 67% of model variance; translated findings into race strategy recommendations reducing pit-stop timing error by 15% vs baseline

### F1 Visual Analytics Dashboard (Tableau)

- Developed 4-dashboard Tableau suite analyzing 100k+ lap records across 10+ seasons uncovering driver performance trends, constructor patterns, and competitive insights
- Designed 15+ calculated KPIs - win rate, YoY improvement, consistency metrics - using LOD expressions and calculated fields for multi-level aggregations
- Implemented interactive drill-down workflows with dashboard actions allowing stakeholders to explore patterns from season-level down to lap-by-lap granularity
- Delivered executive-ready data storytelling translating complex multi-season patterns into actionable insights for non-technical stakeholders

### NYC Yellow Taxi - Cloud Data Warehouse & BI Analytics

- Architected end-to-end ELT pipeline on Snowflake ingesting 2.96M NYC Yellow Taxi trip records (50MB Parquet) using a 3-layer architecture (Raw → Staging → Star Schema) for scalable analytical querying
- Designed Kimball-style dimensional model with 1 fact table and 4 dimension tables across 262 pickup zones, 5 payment types, and 7 rate codes - supporting multi-dimensional revenue and operational analysis
- Engineered 10+ SQL transformations in the staging layer - timestamp conversion, null handling, trip duration derivation, and data quality filters - removing 92K noisy rows (~3% of dataset)
- Uncovered \$78.5M in January 2024 revenue across 2.87M trips; surfaced insights on peak demand (6–8PM), payment patterns (80% credit card), and top revenue zones (Midtown/Upper East Side)
- Delivered interactive Tableau dashboard with dynamic parameter controls and DAY\_TYPE filters across 6 charts, enabling self-service BI for stakeholder decision-making

## WORK EXPERIENCE

### Coditas Solutions LLP

Aug 2022 - Dec 2024

#### Software Engineer | E-commerce & Healthcare Data Operations

Pune, India

- Streamlined payment reconciliation processing 50k+ financial records daily using SQL and Java, reducing data discrepancies by ~60% through validation logic and anomaly detection
- Deployed healthcare reporting platform serving 1k+ users with 15+ data APIs for patient analytics, ensuring HIPAA compliance and PHI data security
- Established data pipelines on GCP (Pub/Sub, Cloud Scheduler) integrating 11+ external APIs, processing 5k–10k transactions daily with real-time data quality monitoring
- Produced client-facing presentations using Tableau and Excel, collaborating with analytics teams to communicate data insights and business recommendations to stakeholders
- Reduced MySQL query response times by 25% through performance tuning, strategic indexing, and schema redesign
- Implemented automated data validation workflows processing 50k+ daily records, achieving 99.5%+ accuracy across multi-channel e-commerce operations
- Partnered with cross-functional stakeholders to translate business requirements into technical solutions, maintaining comprehensive documentation and 80%+ code coverage

## TECHNICAL SKILLS

- **Machine Learning:** Scikit-learn, Random Forest, Gradient Boosting, CatBoost, cross-validation, hyperparameter tuning, feature engineering, ensemble methods
- **Data Analysis:** Python (Pandas, NumPy, Matplotlib, Seaborn), R, SQL, Excel, EDA, statistical analysis, anomaly detection
- **Data Visualization:** Tableau, PowerBI, Matplotlib, Seaborn, dashboard design, KPI development, LOD expressions, data storytelling, executive presentations
- **Database & SQL:** MySQL, PostgreSQL, Snowflake, Azure Data Studio, window functions, CTEs, aggregate functions, stored procedures, query optimization, indexing
- **Data Engineering:** ELT/ETL pipelines, star schema, Parquet, data warehousing, data transformation, data quality, data validation, data modeling
- **Cloud & Big Data:** GCP (Pub/Sub, Cloud Scheduler, App Engine), Snowflake, AWS, Azure, big data concepts
- **Development:** Python, Java, SQL, Django, Git/GitHub, Jupyter Notebooks, REST APIs, Agile/Scrum
- **Domain Knowledge:** Healthcare data (HIPAA), e-commerce analytics, financial data, operations analytics

## CERTIFICATIONS

- **Clean Code Championship (CCC):** Coditas Technologies | Aug 2023
- **Certified in Java Backend Development & Clean Coding:** Coditas Technologies | Dec 2022